

# PROBLEM STATEMENT

You are a Data scientist at the Forest bank, and the bank has been facing a number of fraudulent transactions. Fraudulent transactions are those which are not done by the user itself but by a hacker or by a card thief or by doing a phishing attack. The bank has a huge dataset of its credit card users and the bank has also done the job of marking these transactions as fraudulent or not. You being banks' Data science expert have been given the task to first analyse this data and generate insights and then predict if a transaction on a credit card is going to be fraud or not.

# IMPORT LIBRARIES

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score
import warnings
warnings.filterwarnings("ignore")
```

# UNDERSTAND DATASET

```
train_data= pd.read_csv("fraudTrain.csv")
test_data=pd.read_csv("fraudTest.csv")
```

## Data Dictionary

- **transdate****trans\_time**: The date and time of the transaction.
- **cc\_num**: credit card number.
- **merchant**: Merchant who was getting paid.
- **category**: In what area does that merchant deal.
- **amt**: Amount of money in American Dollars.

- **first:** first name of the card holder.
- **last:** last name of the card holder.
- **gender:** Gender of the cardholder. Just male and female!
- **street:** Street of card holder residence
- **city:** city of card holder residence
- **state:** state of card holder residence
- **zip:** ZIP code of card holder residence
- **lat:** latitude of card holder
- **long:** longitude of card holder
- **city\_pop:** Population of the city
- **job:** trade of the card holder
- **dob:** Date of birth of the card holder
- **trans\_num:** Transaction ID
- **unix\_time:** Unix time which is the time calculated since 1970 to today.
- **merch\_lat:** latitude of the merchant
- **merch\_long:** longitude of the merchant
- **is\_fraud:** Whether the transaction is fraud(1) or not(0)

```
train_data.head()
```

	Unnamed: 0	trans_date	trans_time	cc_num	\
0	0	2019-01-01	00:00:18	2703186189652095	
1	1	2019-01-01	00:00:44	630423337322	
2	2	2019-01-01	00:00:51	38859492057661	
3	3	2019-01-01	00:01:16	3534093764340240	
4	4	2019-01-01	00:03:06	375534208663984	

	first	\	merchant	category	amt
0	Jennifer		fraud_Rippin, Kub and Mann	misc_net	4.97
1	Stephanie		fraud_Heller, Gutmann and Zieme	grocery_pos	107.23
2	Edward		fraud_Lind-Buckridge	entertainment	220.11
3			fraud_Kutch, Hermiston and Farrell	gas_transport	45.00

```

Jeremy
4          fraud_Keeling-Crist      misc_pos      41.96
Tyler

      last gender      street ...      lat
long \
0   Banks      F      561 Perry Cove ... 36.0788 -
81.1781
1   Gill      F  43039 Riley Greens Suite 393 ... 48.8878 -
118.2105
2   Sanchez      M      594 White Dale Suite 530 ... 42.1808 -
112.2620
3   White      M  9443 Cynthia Court Apt. 038 ... 46.2306 -
112.1138
4   Garcia      M      408 Bradley Rest ... 38.4207 -
79.4629

```

```

      city_pop      job      dob \
0      3495      Psychologist, counselling 1988-03-09
1      149      Special educational needs teacher 1978-06-21
2      4154      Nature conservation officer 1962-01-19
3      1939      Patent attorney 1967-01-12
4      99      Dance movement psychotherapist 1986-03-28

```

```

      trans_num      unix_time      merch_lat      merch_long
\
0  0b242abb623afc578575680df30655b9  1325376018  36.011293  -82.048315
1  1f76529f8574734946361c461b024d99  1325376044  49.159047  -118.186462
2  a1a22d70485983eac12b5b88dad1cf95  1325376051  43.150704  -112.154481
3  6b849c168bdad6f867558c3793159a81  1325376076  47.034331  -112.561071
4  a41d7549acf90789359a9aa5346dcb46  1325376186  38.674999  -78.632459

```

```

      is_fraud
0      0
1      0
2      0
3      0
4      0

```

[5 rows x 23 columns]

```
test_data.head()
```

```

      Unnamed: 0      trans_date      trans_time      cc_num \
0      0      2020-06-21 12:14:25  2291163933867244
1      1      2020-06-21 12:14:33  3573030041201292

```

2	2	2020-06-21 12:14:53	3598215285024754
3	3	2020-06-21 12:15:15	3591919803438423
4	4	2020-06-21 12:15:17	3526826139003047

	merchant	category	amt	first
\				
0	fraud_Kirlin and Sons	personal_care	2.86	Jeff
1	fraud_Sporer-Keebler	personal_care	29.84	Joanne
2	fraud_Swaniawski, Nitzsche and Welch	health_fitness	41.28	Ashley
3	fraud_Haley Group	misc_pos	60.05	Brian
4	fraud_Johnston-Casper	travel	3.19	Nathan

	last	gender	street	...	lat
long \					
0	Elliott	M	351 Darlene Green	...	33.9659 -
80.9355					
1	Williams	F	3638 Marsh Union	...	40.3207 -
110.4360					
2	Lopez	F	9333 Valentine Point	...	40.6729 -
73.5365					
3	Williams	M	32941 Krystal Mill Apt. 552	...	28.5697 -
80.8191					
4	Massey	M	5783 Evan Roads Apt. 465	...	44.2529 -
85.0170					

	city_pop	job	dob	\
0	333497	Mechanical engineer	1968-03-19	
1	302	Sales professional, IT	1990-01-17	
2	34496	Librarian, public	1970-10-21	
3	54767	Set designer	1987-07-25	
4	1126	Furniture designer	1955-07-06	

	trans_num	unix_time	merch_lat	merch_long
\				
0	2da90c7d74bd46a0caf3777415b3ebd3	1371816865	33.986391	-81.200714
1	324cc204407e99f51b0d6ca0055005e7	1371816873	39.450498	-109.960431
2	c81755dbbba9d5c77f094348a7579be	1371816893	40.495810	-74.196111
3	2159175b9efe66dc301f149d3d5abf8c	1371816915	28.812398	-80.883061
4	57ff021bd3f328f8738bb535c302a31b	1371816917	44.959148	-85.884734

is\_fraud

```
0      0
1      0
2      0
3      0
4      0
```

```
[5 rows x 23 columns]
```

```
train_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 1296675 entries, 0 to 1296674
```

```
Data columns (total 23 columns):
```

#	Column	Non-Null Count	Dtype
0	Unnamed: 0	1296675 non-null	int64
1	trans_date_trans_time	1296675 non-null	object
2	cc_num	1296675 non-null	int64
3	merchant	1296675 non-null	object
4	category	1296675 non-null	object
5	amt	1296675 non-null	float64
6	first	1296675 non-null	object
7	last	1296675 non-null	object
8	gender	1296675 non-null	object
9	street	1296675 non-null	object
10	city	1296675 non-null	object
11	state	1296675 non-null	object
12	zip	1296675 non-null	int64
13	lat	1296675 non-null	float64
14	long	1296675 non-null	float64
15	city_pop	1296675 non-null	int64
16	job	1296675 non-null	object
17	dob	1296675 non-null	object
18	trans_num	1296675 non-null	object
19	unix_time	1296675 non-null	int64
20	merch_lat	1296675 non-null	float64
21	merch_long	1296675 non-null	float64
22	is_fraud	1296675 non-null	int64

```
dtypes: float64(5), int64(6), object(12)
```

```
memory usage: 227.5+ MB
```

```
train_data.columns
```

```
Index(['Unnamed: 0', 'trans_date_trans_time', 'cc_num', 'merchant',  
      'category', 'amt', 'first', 'last', 'gender', 'street', 'city', 'state',  
      'zip', 'lat', 'long', 'city_pop', 'job', 'dob', 'trans_num',  
      'unix_time',
```

```
'merch_lat', 'merch_long', 'is_fraud'],
dtype='object')
```

## Data Cleaning

Data cleaning by removing all unimportant columns that contains personal information and missing values

```
columns_to_drop = ["Unnamed: 0", "trans_date_trans_time", "cc_num", "first", "last", 'street', 'zip', 'lat', 'long', 'dob', 'trans_num', 'unix_time', 'merch_lat', 'merch_long']
train_data = train_data.drop(columns=columns_to_drop)
# 'category', 'gender', 'job', 'merchant', 'amt', 'city', 'city_pop', 'state', 'is_fraud'
# columns_to_drop = ["Unnamed: 0", "trans_date_trans_time", "cc_num", "first", "last", 'street', 'zip', 'lat', 'long', 'dob', 'trans_num', 'unix_time', 'merch_lat', 'merch_long']
# columns_to_drop = ['trans_date_trans_time', "cc_num", "first", "last"]

columns_to_drop = ["Unnamed: 0", "trans_date_trans_time", "cc_num", "first", "last", 'street', 'zip', 'lat', 'long', 'dob', 'trans_num', 'unix_time', 'merch_lat', 'merch_long']
test_data = test_data.drop(columns=columns_to_drop)

train_data.head()
```

	merchant	category	amt	gender	\
0	fraud_Rippin, Kub and Mann	misc_net	4.97	F	
1	fraud_Heller, Gutmann and Zieme	grocery_pos	107.23	F	
2	fraud_Lind-Buckridge	entertainment	220.11	M	
3	fraud_Kutch, Hermiston and Farrell	gas_transport	45.00	M	
4	fraud_Keeling-Crist	misc_pos	41.96	M	

	city	state	city_pop	job
is_fraud				
0	Moravian Falls	NC	3495	Psychologist, counselling
0				
1	Orient	WA	149	Special educational needs teacher
0				
2	Malad City	ID	4154	Nature conservation officer
0				
3	Boulder	MT	1939	Patent attorney
0				
4	Doe Hill	VA	99	Dance movement psychotherapist
0				

```
test_data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 555719 entries, 0 to 555718
Data columns (total 9 columns):
#   Column      Non-Null Count  Dtype
---  -
0   merchant    555719 non-null    object
1   category    555719 non-null    object
2   amt         555719 non-null    float64
3   gender      555719 non-null    object
4   city        555719 non-null    object
5   state       555719 non-null    object
6   city_pop    555719 non-null    int64
7   job         555719 non-null    object
8   is_fraud    555719 non-null    int64
dtypes: float64(1), int64(2), object(6)
memory usage: 38.2+ MB

train_data.shape

(1296675, 9)

```

There are overall total 1296675 rows and 9 columns in train dataset

```

test_data.shape

(555719, 9)

```

There are overall total 555719 rows and 9 columns in train dataset.

```

train_data.isna().sum()

merchant      0
category      0
amt           0
gender        0
city          0
state         0
city_pop      0
job           0
is_fraud      0
dtype: int64

test_data.isna().sum()

merchant      0
category      0
amt           0
gender        0
city          0
state         0

```

```
city_pop    0
job         0
is_fraud    0
dtype: int64
```

There are no null values present in dataset. Hence, there is no need to perform missing value treatment.


```
train_data.columns
Index(['merchant', 'category', 'amt', 'gender', 'city', 'state',
      'city_pop',
      'job', 'is_fraud'],
      dtype='object')
```

## EXPLORATORY DATA ANALYSIS


```
plt.figure(figsize=(15,2))
sns.pairplot(train_data[['category', 'gender', 'job', 'merchant',
                        'amt', 'city', 'city_pop', 'state', 'is_fraud']], hue='is_fraud')
plt.show()
```

<Figure size 1500x200 with 0 Axes>






```
plt.figure(figsize=(20,6))
sns.countplot(data=train_data, x = 'gender',hue='is_fraud')
plt.show()
```



```
train_data[train_data.is_fraud==1]['amt'].hist()
<Axes: >
```



As can be seen in the histogram, around 1750+ fraudulent transactions were in the bucket 0-100 dollars and least on the 1200-1400 dollars bucket

## ENCODING

```
numerical = []
categorical = []
for i in train_data.columns:
    if train_data[i].dtype == 'int64' or train_data[i].dtype == 'float64':
        numerical.append(i)
    else:
        categorical.append(i)

numerical_test = []
categorical_test = []
for i in test_data.columns:
    if test_data[i].dtype == 'int64' or test_data[i].dtype == 'float64':
        numerical_test.append(i)
    else:
        categorical_test.append(i)
```



```
train_data.head()
```

	merchant	category	amt	gender	city	state	city_pop	job
is_fraud								
0	514	8	-0.407826	0	526	27	-0.282589	370
0								
1	241	4	0.230039	0	612	47	-0.293670	428
0								
2	390	0	0.934149	1	468	13	-0.280406	307
0								
3	360	2	-0.158132	1	84	26	-0.287742	328
0								
4	297	9	-0.177094	1	216	45	-0.293835	116
0								

## CLASSIFICATIONS MODEL TO PREDICT FRAUDLENT OR NOT

```
x_train = train_data.drop('is_fraud', axis=1)
y_train = train_data['is_fraud']
```

```
x_test = test_data.drop('is_fraud', axis=1)
y_test = test_data['is_fraud']
```

## MODEL 1: LOGISTIC REGRESSION

```
logistic_reg = LogisticRegression()
logistic_reg.fit(x_train, y_train)
y_pred_logistic = logistic_reg.predict(x_test)
accuracy_logistic = accuracy_score(y_test, y_pred_logistic)
print(f"Logistic Regression Accuracy: {accuracy_logistic}")
```

Logistic Regression Accuracy: 0.9955139198047934

## MODEL 2:KNN

```
knn = KNeighborsClassifier()
knn.fit(x_train, y_train)
y_pred_knn = knn.predict(x_test)
accuracy_knn = accuracy_score(y_test, y_pred_knn)
print(f"K-Nearest Neighbors Accuracy: {accuracy_knn}")
```

K-Nearest Neighbors Accuracy: 0.9958540197473903


```

from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
cm= confusion_matrix(y_test, y_pred_knn)

cm_disp= ConfusionMatrixDisplay(confusion_matrix=cm)
cm_disp.plot()

<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at
0x1ef890fac10>

```



```

from sklearn.metrics import classification_report, accuracy_score,
log_loss

target_names= ["Negative(0)","Positive(1)"]
# Classification Report
print(classification_report(y_test,knn.predict(x_test),target_names=ta
rget_names))

```

	precision	recall	f1-score	support
Negative(0)	1.00	1.00	1.00	553574
Positive(1)	0.02	0.00	0.00	2145
accuracy			1.00	555719
macro avg	0.51	0.50	0.50	555719

weighted avg	0.99	1.00	0.99	555719
--------------	------	------	------	--------

```
accuracy_score(y_test, y_pred_knn)
```

```
0.9958540197473903
```

## MODEL 3: DECISION TREE

```
decision_tree = DecisionTreeClassifier()  
decision_tree.fit(x_train, y_train)  
y_pred_tree = decision_tree.predict(x_test)  
accuracy_tree = accuracy_score(y_test, y_pred_tree)  
print(f"Decision Tree Accuracy: {accuracy_tree}")
```


```
Decision Tree Accuracy: 0.9958810118063266
```

```
cm_dt= confusion_matrix(y_test, y_pred_tree)
```

```
cm_disp= ConfusionMatrixDisplay(confusion_matrix=cm_dt)
```

```
cm_disp.plot()
```

```
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at  
0x1ef88fa64d0>
```



```
target_names= ["Negative(0)","Positive(1)"]
# Classification Report
print(classification_report(y_test,decision_tree.predict(x_test),target_names=target_names))
```

	precision	recall	f1-score	support
Negative(0)	1.00	1.00	1.00	553574
Positive(1)	0.47	0.59	0.52	2145
accuracy			1.00	555719
macro avg	0.74	0.79	0.76	555719
weighted avg	1.00	1.00	1.00	555719

```
accuracy_score(y_test, y_pred_tree)
```

```
0.9958810118063266
```

## MODEL 4: RANDOM FOREST

```
random_forest = RandomForestClassifier()
random_forest.fit(x_train, y_train)
```

```


y_pred_forest = random_forest.predict(x_test)
accuracy_forest = accuracy_score(y_test, y_pred_forest)
print(f"Random Forest Accuracy: {accuracy_forest}")

Random Forest Accuracy: 0.9976444929901623

cm_rf= confusion_matrix(y_test, y_pred_forest)
cm_disp= ConfusionMatrixDisplay(confusion_matrix=cm_rf)
cm_disp.plot()

<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at
0x1ef891fb9d0>

```



```

target_names= ["Negative(0)", "Positive(1)"]
# Classification Report
print(classification_report(y_test, random_forest.predict(x_test), target_names=target_names))

```

	precision	recall	f1-score	support
Negative(0)	1.00	1.00	1.00	553574
Positive(1)	0.75	0.58	0.66	2145



accuracy			1.00	555719
macro avg	0.88	0.79	0.83	555719
weighted avg	1.00	1.00	1.00	555719

```
accuracy_score(y_test, y_pred_forest)
```

```
0.9976444929901623
```

## MODEL 5: SVM

```
svm = SVC()
svm.fit(x_train, y_train)
y_pred_svm = svm.predict(x_test)
accuracy_svm = accuracy_score(y_test, y_pred_svm)
print(f"SVM Accuracy: {accuracy_svm}")
```


```
SVM Accuracy: 0.9961401355721147
```

```
cm_svm= confusion_matrix(y_test, y_pred_svm)
```

```
cm_disp= ConfusionMatrixDisplay(confusion_matrix=cm_svm)
```

```
cm_disp.plot()
```

```
<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at  
0x1ef9a77c810>
```



```
target_names= ["Negative(0)","Positive(1)"]
# Classification Report
print(classification_report(y_test,svm.predict(x_test),target_names=target_names))
```

	precision	recall	f1-score	support
Negative(0)	1.00	1.00	1.00	553574
Positive(1)	0.00	0.00	0.00	2145
accuracy			1.00	555719
macro avg	0.50	0.50	0.50	555719
weighted avg	0.99	1.00	0.99	555719

```
accuracy_score(y_test, y_pred_svm)
```

```
0.9961401355721147
```

All the 5 Models are giving excellent results with more than 99% of accuracy and hence there is no need of tuning any parameter.